

## Table B — brainmed-8b-v1\_\_soup-a0.3 vs Huatuo-o1-RL-8B

Accuracy (%). Published values from MedReason Table 4; our column measured by this pipeline.

| Benchmark | Huatuo-o1-RL-8B (published) | brainmed-8b-v1__soup-a0.3 (measured) | Delta |     |  |
|-----------|-----------------------------|--------------------------------------|-------|-----|--|
| MB-op4    | 55.20                       | 59.42                                | +4.22 | win |  |
| MB-op5    | 51.30                       | 54.87                                | +3.57 | win |  |
| MedXpert  | 16.70                       | 16.43                                | -0.27 |     |  |
| MedQA     | 72.60                       | 76.28                                | +3.68 | win |  |
| MedMCQA   | 60.40                       | 62.54                                | +2.14 | win |  |
| PubMedQA  | 79.20                       | 80.00                                | +0.80 | win |  |
| Average   | 55.90                       | 58.26                                | +2.36 | win |  |

Excluded from this table: MedReason-8B.

Backbone shown with its published values (Huatuo-o1-RL-8B), not the score measured on this harness.

Harness offset measured on the untouched backbone: +2.89 pt vs the published protocol. Subtract it before reading our-vs-published gaps.

Wins on 5 of 6 benchmarks.