

Table D — Eight-benchmark suite (MedReason Table 2 columns)

Accuracy (%). Published rows use different backbones - compare within a backbone family, not across.

Model / data	MedQA	MedMCQA	PubMedQA	MMLU-Pro	MB-op4	MB-op5	MedXpert	HLE(med)	Avg
Llama3.1-Instruct-8B (base)	58.70	56.00	75.20	58.20	48.70	42.50	13.20	13.60	45.80
Llama3.1-Instruct-8B + Huatuo CoT	70.20	58.20	76.10	59.90	53.30	49.70	17.30	14.60	49.90
Llama3.1-Instruct-8B + MedReason	68.40	57.50	77.60	63.10	57.50	52.30	16.40	16.50	51.20
Mistral-Instruct-7B (base)	48.20	44.90	50.10	42.70	43.50	33.40	11.40	14.60	36.10
Mistral-Instruct-7B + Huatuo CoT	59.90	46.90	57.50	47.60	50.00	46.10	14.40	14.60	42.10
Mistral-Instruct-7B + MedReason	58.70	48.90	59.20	50.80	52.30	47.10	16.60	24.30	44.70
brainmed-8b-final (ours)	76.67	64.07	79.10	65.73	62.01	54.87	18.63	10.68	53.97
brainmed-8b-v1 (ours)	72.58	60.22	77.70	61.82	56.49	50.65	17.81	20.39	52.21

Excluded from this table: MedReason-8B.

Backbone shown with its published values (Huatuo-o1-RL-8B), not the score measured on this harness.

Harness offset measured on the untouched backbone: +2.89 pt vs the published protocol. Subtract it before reading our-vs-published gaps.